

# **AI-based autonomous agents - What, why and how?**

07.04.2025

**Presenter: Dr. Iurii Savvateev**

# Outline

Part I. What are AI-based autonomous agents?

Break I

Part II. Why and in which areas to use agents?

Break II

Part III. How to practically start using agents?

# Outline

## Part I. What are AI-based autonomous agents?

Break I

Part II. Why and in which areas to use agents?

Break II

Part III. How to practically start using agents?

# Part I. What?

We define AI as the study of agents that receive percepts from the environment and perform actions.


---

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



- Perception
- Action
- Decision-making
- Autonomy



Any entity capable of perceiving its environment and taking actions can be considered an agent  
(Cheng et al., 2024)

Images were created with biorender.com

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



(Russel and Norvig, 2010; Cheng et al., 2024)

<https://www.digitalocean.com/resources/articles/types-of-ai-agents>

Images were created with biorender.com

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

Agent



Actions

Percepts



Actions based solely  
on current percepts

Model-based Reflex  
Agents

Goal-based Agents

Utility-based Agents

Learning Agents

(Russel and Norvig, 2010; Cheng et al., 2024)

<https://www.digitalocean.com/resources/articles/types-of-ai-agents>

Images were created with biorender.com

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

Agent



Actions

Percepts



Actions based solely  
on current percepts

Environment model +  
condition-action rules

Goal-based Agents

Utility-based Agents

Learning Agents

(Russel and Norvig, 2010; Cheng et al., 2024)

<https://www.digitalocean.com/resources/articles/types-of-ai-agents>


Images were created with biorender.com



# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Actions based solely  
on current percepts

Environment model +  
condition-action rules

Environment model +  
goal-based state  
evaluation

Utility-based Agents

Learning Agents

(Russel and Norvig, 2010; Cheng et al., 2024)

<https://www.digitalocean.com/resources/articles/types-of-ai-agents>

Images were created with biorender.com

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

Agent



Actions

Percepts



Actions based solely  
on current percepts

Environment model +  
condition-action rules

Environment model +  
goal-based state  
evaluation

Environment model +  
utility function for  
state evaluation with  
trade offs

Learning Agents

(Russel and Norvig, 2010; Cheng et al., 2024)


<https://www.digitalocean.com/resources/articles/types-of-ai-agents>

Images were created with biorender.com

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Simple Reflex Agents

Model-based Reflex  
Agents

Goal-based Agents

Utility-based Agents

Reinforcement  
Learning based  
agents

Large Language  
Model based agents

(Russel and Norvig, 2010; Cheng et al., 2024)


<https://www.digitalocean.com/resources/articles/types-of-ai-agents>

Images were created with biorender.com

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Capacity to learn and improve based on experience



Autonomous learning capabilities and multi-step decision problems

(Russel and Norvig, 2010; Cheng et al., 2024)


<https://www.digitalocean.com/resources/articles/types-of-ai-agents>

Images were created with biorender.com

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



## Chess




Silver et al. 2018

Learn unknown policy to maximize cumulative rewards



Teslamag, 2017




DeepMind, 2019

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Simple Reflex Agents

Model-based Reflex Agents

Goal-based Agents

Utility-based Agents

Reinforcement Learning based agents

Large Language Model based agents

Capacity to learn and improve based on experience




Autonomous learning capabilities and multi-step decision problems

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Simple Reflex Agents

Model-based Reflex  
Agents

Goal-based Agents

Utility-based Agents

Reinforcement  
Learning based  
agents

Large Language  
Model based agents


## Limitations:

- Long training time
- Sample efficiency (high number of samples needed)
- Instability in non-stationary environments (change of the dynamics as policy evolves)
- Lack of generalizability (partially tackled with transfer learning, but still an active research area)

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Simple Reflex Agents

Model-based Reflex  
Agents

Goal-based Agents

Utility-based Agents

Reinforcement  
Learning based  
agents


Large Language  
Model based agents



# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



| Model Name        | Developer       | Country of Origin |
|-------------------|-----------------|-------------------|
| GPT-4.5           | OpenAI          | United States     |
| DeepSeek R1       | DeepSeek        | China             |
| Claude 3.7 Sonnet | Anthropic       | United States     |
| Grok-3            | xAI             | United States     |
| Mistral Large 2   | Mistral AI      | France            |
| PaLM 2            | Google          | United States     |
| Gemini 2.0 Pro    | Google DeepMind | United States     |
| Llama 4 Maverick  | Meta            | United States     |
| Command R+        | Cohere          | Canada            |
| ERNIE 4.0         | Baidu           | China             |

Language Overview Price Analysis WebDev Arena Vision Text-to-Image Copilot Arena Arena-Hard-Auto

Total #models: 222. Total #votes: 2,838,248. Last updated: 2025-04-02.

Code to recreate leaderboard tables and plots in this [notebook](#). You can contribute your vote at [lmarena.ai](#)!

Category: Overall Apply filter: ☐ Style Control ☐ Show Deprecated

Overall Questions  
#models: 222 (100%) #votes: 2,838,248 (100%)


| Rank* (UB) | Rank (StyleCtrl) | Model                                               | Arena Score | 95% CI  | Votes | Organization | License     |
|------------|------------------|-----------------------------------------------------|-------------|---------|-------|--------------|-------------|
| 1          | 1                | <a href="#">Gemini-2.5-Pro-Exp-03-25</a>            | 1439        | +7/-10  | 5858  | Google       | Proprietary |
| 2          | 5                | <a href="#">Llama-4-Maverick-03-26-Experimental</a> | 1417        | +13/-12 | 2520  | Meta         | Llama       |
| 2          | 1                | <a href="#">ChatGPT-4o-latest-(2025-03-26)</a>      | 1410        | +8/-10  | 4899  | OpenAI       | Proprietary |
| 2          | 4                | <a href="#">Grok-3-Preview-02-24</a>                | 1403        | +6/-6   | 12391 | xAI          | Proprietary |
| 3          | 2                | <a href="#">GPT-4.5-Preview</a>                     | 1398        | +5/-7   | 12312 | OpenAI       | Proprietary |
| 6          | 7                | <a href="#">Gemini-2.0-Flash-Thinking-Exp-01-21</a> | 1380        | +4/-4   | 24298 | Google       | Proprietary |
| 6          | 4                | <a href="#">Gemini-2.0-Pro-Exp-02-05</a>            | 1380        | +4/-4   | 20289 | Google       | Proprietary |
| 6          | 4                | <a href="#">DeepSeek-V3-0324</a>                    | 1369        | +10/-10 | 3526  | DeepSeek     | MIT         |
| 8          | 5                | <a href="#">DeepSeek-R1</a>                         | 1358        | +5/-5   | 14259 | DeepSeek     | MIT         |

<https://lmarena.ai/?leaderboard>

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



| Model Name        | Developer       | Country of Origin |
|-------------------|-----------------|-------------------|
| GPT-4.5           | OpenAI          | United States     |
| DeepSeek R1       | DeepSeek        | China             |
| Claude 3.7 Sonnet | Anthropic       | United States     |
| Grok-3            | xAI             | United States     |
| Mistral Large 2   | Mistral AI      | France            |
| PaLM 2            | Google          | United States     |
| Gemini 2.0 Pro    | Google DeepMind | United States     |
| Llama 4 Maverick  | Meta            | United States     |
| Command R+        | Cohere          | Canada            |
| ERNIE 4.0         | Baidu           | China             |


| Model                               | Organization | Global Average | Reasoning Average | Coding Average | Mathematics Average | Data Analysis Average | Language Average |
|-------------------------------------|--------------|----------------|-------------------|----------------|---------------------|-----------------------|------------------|
| gemini-2.5-pro-exp-03-25            | Google       | 82.71          | 89.75             | 85.87          | 90.20               | 79.89                 | 67.82            |
| deepseek-r1                         | DeepSeek     | 69.79          | 83.17             | 66.74          | 80.71               | 69.78                 | 48.53            |
| o1-2024-12-17-high                  | OpenAI       | 74.49          | 91.58             | 69.69          | 80.32               | 65.47                 | 65.39            |
| claude-3.7-sonnet-thinking          | Anthropic    | 75.07          | 87.83             | 74.54          | 79.00               | 74.05                 | 59.93            |
| qwq-32b                             | Alibaba      | 69.99          | 83.50             | 72.23          | 77.82               | 65.03                 | 51.35            |
| o3-mini-2025-01-31-high             | OpenAI       | 74.19          | 89.58             | 82.74          | 77.29               | 70.64                 | 50.68            |
| gemini-2.0-flash-thinking-exp-01-21 | Google       | 63.81          | 78.17             | 53.49          | 75.85               | 69.37                 | 42.18            |
| deepseek-v3-0324                    | DeepSeek     | 63.94          | 65.83             | 70.90          | 73.52               | 60.37                 | 49.07            |
| o3-mini-2025-01-31-medium           | OpenAI       | 67.38          | 86.33             | 65.38          | 72.37               | 66.56                 | 46.26            |
| gemini-exp-1206                     | Google       | 61.44          | 57.00             | 63.41          | 72.36               | 63.16                 | 51.29            |
| gemini-2.0-pro-exp-02-05            | Google       | 61.48          | 60.08             | 63.49          | 70.97               | 68.02                 | 44.85            |

<https://livebench.ai/#/>

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.


— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).




## Attention is all you need (Vaswani et al., 2017)

- Text => Embedding
- **Attention block** captures the context: Embedding of Part 1  $\Leftrightarrow$  Embedding of Part 2 etc.
- **Feed forward block** links the data with the stored facts: Embedding  $\Leftrightarrow$  facts
- Embedding => Text

Scaled Dot-Product Attention




Multi-Head Attention



# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



## Attention is all you need (Vaswani et al., 2017)

- Text => Embedding
- **Attention block** captures the context: Embedding of Part 1  $\Leftrightarrow$  Embedding of Part 2 etc.
- **Feed forward block** links the data with the stored facts: Embedding  $\Leftrightarrow$  facts
- Embedding => Text




Figure 1: The Transformer - model architecture.

# Part I. Definitions


We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

$$\frac{e^{x_i/T}}{\sum_{n=0}^{N-1} e^{x_n/T}}$$

- Text => Embedding
- **Attention block** captures the context:  
Embedding of Part 1 ⇔ Embedding of Part 2 etc.
- **Feed forward block** links the data with the stored facts: Embedding ⇔ facts
- Embedding => Text


Figure from Hongying et al.2025



# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



N-shots prompting (N refers to the example numbers)

## Model input

Q: Yes or no: is it common to see frost during some college commencements?  
A: The answer is yes.  
Q: Yes or no: would a pear sink in water?

## Model output


A: The answer is No. ✓

Figure from Chen et al. 2023

# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



## Model input

Imagine an infinitely wide entrance, which is more likely to pass through it, a military tank or a car? Let's think step by step.

## Model output

Given an infinitely wide entrance, both a military tank and a car could pass through without any issues related to the width. However, there are still several factors that could influence which is more likely to pass through:

1. **Ground surface:** tanks are designed to handle rough terrains that cars may struggle with. If the ground leading to and from the entrance is too rugged or steep for a car, the tank would be more likely to pass through.
2. **Weight restrictions:** if the entrance has a load limit due to the materials used in its construction, the car, being lighter, could be more likely to pass through. The average car weighs around 1.5 tons, while a military tank can weigh anywhere from 60 to 70 tons.
3. **Height clearance:** if the entrance is infinitely wide but not infinitely tall, the height of the vehicle matters. The average height of a car is around 1.5 meters, while tanks can be around 2.4 meters. If the entrance is shorter than around 2.4 meters, a car would be more likely to pass through.
4. **Regulations and permissions:**

...  
In conclusion, without specific information about the conditions at the entrance, we would assume that a car is more likely to pass through, simply because cars are more common, more likely to meet potential height and weight restrictions and face fewer regulatory barriers.


N-shots prompting (N refers to the example numbers)  
Chain of thoughts

Figure from Chen et al. 2023


# Part I. Definitions

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003)




N-shots prompting (N refers to the example numbers)  
Chain of thoughts



Figures are from Kojima et al., 2022



# Part I. LLM-based agents



We define AI as the study of agents that receive percepts from the environment and perform actions.


— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

Limitations:

- Context Length Constraint


Figure from

<https://www.artfish.ai/p/long-context-llms>



artfish.ai

# Part I. LLM-based agents



We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

## Limitations:

- Context Length Constraint

### Lost in the Middle: How Language Models Use Long Context:

Nelson F. Liu<sup>1\*</sup> Kevin Lin<sup>2</sup> John Hewitt<sup>1</sup> Ashwin Paranajpe<sup>3</sup>  
Michele Bevilacqua<sup>3</sup> Fabio Petroni<sup>3</sup> Percy Liang<sup>1</sup>  
<sup>1</sup>Stanford University <sup>2</sup>University of California, Berkeley <sup>3</sup>Samaya AI  
nfliu@cs.stanford.edu





Figure from Liu et al., 2023



Images were created with biorender.com

# Part I. LLM-based agents

We define AI as the study of agents that receive percepts from the environment and perform actions.


— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

Limitations:

- Context Length Constraint


Lost in the Middle: How Language Models Use Long Context:

Nelson F. Liu<sup>1\*</sup> Kevin Lin<sup>2</sup> John Hewitt<sup>1</sup> Ashwin Paranajpe<sup>3</sup>  
Michele Bevilacqua<sup>3</sup> Fabio Petroni<sup>3</sup> Percy Liang<sup>1</sup>  
<sup>1</sup>Stanford University <sup>2</sup>University of California, Berkeley <sup>3</sup>Samaya AI  
nflui@cs.stanford.edu



RAG overview figure from


<https://python.langchain.com/v0.2/docs/concepts/>



# Part I. LLM-based agents

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).




## Limitations:

- Context Length Constraint
- Protracted Knowledge Update

## Carbon footprint comparison


Source: Strubell et al, 2019.



Reconstructed from: <http://arxiv.org/abs/1906.02243>

## The cost of the computational power required to train the most powerful AI systems has doubled every nine months

Cost of computational power required to train frontier AI systems



Cost includes amortized hardware acquisition and energy consumption. Red shaded area indicates 95% confidence prediction interval.

Chart: Will Henshall for TIME • Source: Epoch AI • [Get the data](#) • Created with [Datawrapper](#)

# Part I. LLM-based agents




We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).


## Limitations:

- Context Length Constraint
- Protracted Knowledge Update
- No direct Tool Utilization (math)



Yuntian Deng  
@yuntiandeng · Follow


Is OpenAI's o1 a good calculator? We tested it on up to 20x20 multiplication—o1 solves up to 9x9 multiplication with decent accuracy, while gpt-4o struggles beyond 4x4. For context, this task is solvable by a small LM using implicit CoT with stepwise internalization. 1/4



8:46 PM · Sep 17, 2024

3.1K Reply Copy link

# Part I. LLM-based agents



We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).

## Limitations:

- Context Length Constraint
- Protracted Knowledge Update
- No direct Tool Utilization (web)



## Seattle (Q5083)

[Item](#) [Discussion](#)

city in and county seat of King County, State of Washington, United States  
Seattle, Washington | Seattle, WA | Duwamps | New York Alki | Emerald City

Provide wikidata identifier for Berlin city in the US. Only use the stored information and no web search.

The Wikidata identifier for **Berlin, a city in the United States** (specifically **Berlin, New Hampshire**, which is the most prominent Berlin in the US) is [Q5083](#).

Let me know if you meant a different Berlin (like Berlin, Wisconsin or Berlin, Maryland), and I can provide that identifier instead.


< 2/2 >



# Part I. LLM-based agents

We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



## Limitations:

- Context Length Constraint
- Protracted Knowledge Update
- No direct Tool Utilization (web)

Provide wikidata identifier for Berlin city in the US. Use web search

In the United States, there are several cities and towns named Berlin, each with its own Wikidata identifier:

- Berlin, New Hampshire: [Q821244](#) Wikidata
- Berlin, Maryland: [Q614184](#) Wikidata
- Berlin, Wisconsin: [Q1569850](#) Wikidata

If you're interested in a specific Berlin, please let me know, and I can provide more detailed information.

Sources



Berlin (Q821244)

Item

Discussion

city in Coos County, New Hampshire, United States


Berlin, New Hampshire | Berlin, NH

# Part I. LLM-based agents



We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Images were created with biorender.com




# Part I. LLM-based agents



We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).




Images were created with biorender.com

# Part I. LLM-based agents



We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).




Images were created with biorender.com

# Part I. LLM-based agents



We define AI as the study of agents that receive percepts from the environment and perform actions.


— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



# Part I. LLM-based agents

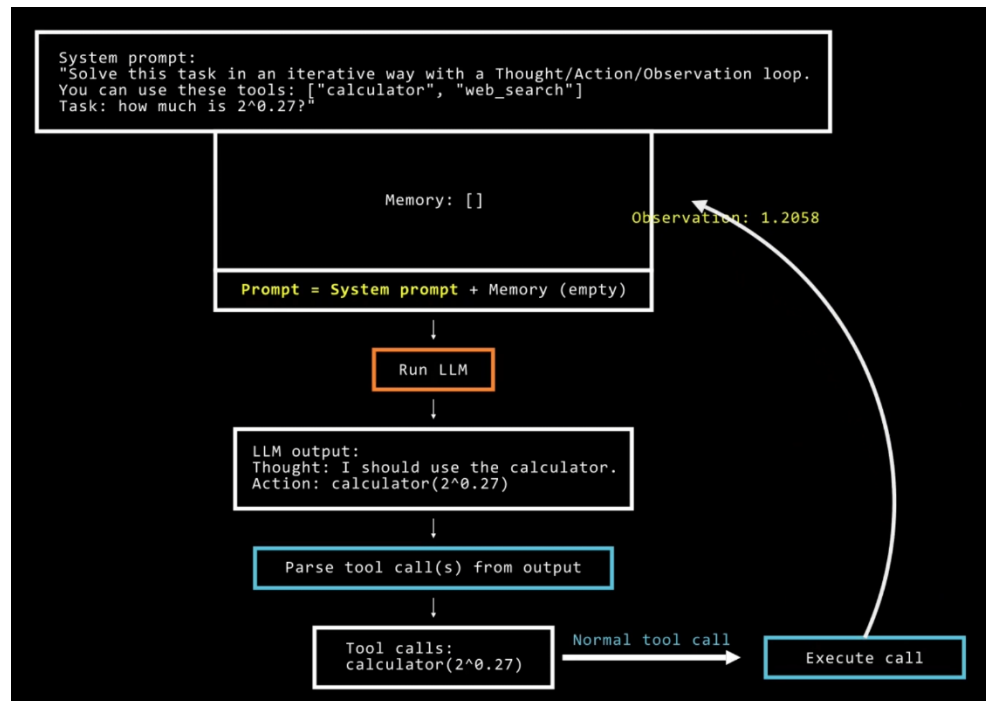
We define AI as the study of agents that receive percepts from the environment and perform actions.

— *Artificial Intelligence: A Modern Approach*, Stuart Russell and Peter Norvig (2003).



Figures are adapted from

[https://huggingface.co/docs/smolagents/en/conceptual\\_guides/intro\\_agents](https://huggingface.co/docs/smolagents/en/conceptual_guides/intro_agents)



# References and useful links:

Stuart J Russell and Peter Norvig. Artificial Intelligence: A Modern Approach (4th Edition). Pearson Education, Inc., 2010.

Cheng, Y., Zhang, C., Zhang, Z., Meng, X., Hong, S., Li, W., Wang, Z., Wang, Z., Yin, F., Zhao, J., & He, X. (2024, January 7). Exploring Large Language Model based Intelligent Agents: Definitions, Methods, and Prospects. arXiv.org. <https://arxiv.org/abs/2401.03428>

7 Types of AI agents to Automate your workflows in 2025 | DigitalOcean. (n.d.). <https://www.digitalocean.com/resources/articles/types-of-ai-agents>

Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., Graepel, T., Lillicrap, T., Simonyan, K., & Hassabis, D. (2018). A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 362(6419), 1140–1144.

<https://doi.org/10.1126/science.aar6404>

<https://deepmind.google/discover/blog/alphazero-shedding-new-light-on-chess-shogi-and-go/>

DeepMind. (2019, January 24). AlphaStar: Mastering the real-time strategy game StarCraft II. <https://deepmind.com/blog/alphastar-mastering-real-time-strategy-game-starcraft-ii/>

Teslamag. (2017, January 16). *Update für den Autopiloten 2.0 wird nun auf alle entsprechenden Fahrzeuge installiert*. Teslamag.de. <https://teslamag.de/news/update-fuer-den-autopiloten-2-0-wird-nun-auf-alle-entsprechenden-fahrzeuge-installiert-11712>

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017, June 12). *Attention is all you need*. arXiv.org. <https://arxiv.org/abs/1706.03762>

[https://en.wikipedia.org/wiki/Attention\\_Is\\_All\\_You\\_Need](https://en.wikipedia.org/wiki/Attention_Is_All_You_Need)

Hongying, Z., Javed, A., Abdullah, M., Rashid, J., & Faheem, M. (2025). Large language models with contrastive decoding algorithm for hallucination mitigation in Low-Resource languages. *CAAI Transactions on Intelligence Technology*. <https://doi.org/10.1049/cit2.70004>

# References and useful links:

Chen, B., Zhang, Z., Langrené, N., & Zhu, S. (2023, October 23). *Unleashing the potential of prompt engineering in Large Language Models: a comprehensive review*. arXiv.org. <https://arxiv.org/abs/2310.14735>

Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., & Iwasawa, Y. (2022, May 24). *Large Language Models are Zero-Shot Reasoners*. arXiv.org. <https://arxiv.org/abs/2205.11916>

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2023, July 6). *Lost in the Middle: How language models use long contexts*. arXiv.org. <https://arxiv.org/abs/2307.03172>

*Q&A with the Experts: Why ChatGPT struggles with math* | Waterloo News. (2024, October 17). Waterloo News. <https://uwaterloo.ca/news/media/qa-experts-why-chatgpt-struggles-math>

Agentic schemas were created with [Biorender.com](https://biorender.com)

# Extra sources for prompt engineering

Tips and Explanations from OpenAI

<https://help.openai.com/en/articles/6654000-best-practices-for-prompt-engineering-with-the-openai-api>

Examples

<https://platform.openai.com/examples>

Prompt engineering tips from HuggingFace

<https://huggingface.co/docs/transformers/en/tasks/prompting>

Prompt engineering tips from Google

<https://ai.google.dev/gemini-api/docs/prompting-intro>

German Prompt Generator

<https://deinkikompass.de/prompt-generator>

# Outline

Part I. What are AI-based autonomous agents?

**Break I**

Part II. Why and in which areas to use agents?

Break II

Part III. How to practically start using agents?



# Outline

Part I. What are AI-based autonomous agents?


Break I

**Part II. Why and in which areas to use agents?**


Break II

Part III. How to practically start using agents?

# Part 2. Single agents



## Part 2. Single Agents



Published as a conference paper at ICLR 2023

### REACT: SYNERGIZING REASONING AND ACTING IN LANGUAGE MODELS

Shunyu Yao<sup>\*1</sup>, Jeffrey Zhao<sup>2</sup>, Dian Yu<sup>2</sup>, Nan Du<sup>2</sup>, Izhak Shafran<sup>2</sup>, Karthik Narasimhan<sup>1</sup>, Yuan Cao<sup>2</sup>

<sup>1</sup>Department of Computer Science, Princeton University



<sup>2</sup>Google Research, Brain team

<sup>1</sup>{shunyuy, karthikn}@princeton.edu


<sup>2</sup>{jeffreyzhao, dianyu, dunan, izhak, yuancao}@google.com

Figure from


<https://www.foundingminds.com/chain-of-thought-reasoning-the-magic-behind-the-o1-model/>



# Part 2. Single agents



# Part 2. Single agents



Published as a conference paper at ICLR 2024

## CRITIC: LARGE LANGUAGE MODELS CAN SELF-CORRECT WITH TOOL-INTERACTIVE CRITIQUING

Zhibin Gou<sup>12\*</sup>, Zhihong Shao<sup>12\*</sup>, Yeyun Gong<sup>2</sup>, Yelong Shen<sup>3</sup>,  
Yujia Yang<sup>1†</sup>, Nan Duan<sup>2</sup>, Weizhu Chen<sup>3</sup>

<sup>1</sup>Tsinghua University

<sup>2</sup>Microsoft Research Asia, <sup>3</sup>Microsoft Azure AI

{gzb22, szh19}@emails.tsinghua.edu.cn, yang.yujia@sz.tsinghua.edu.cn  
{yegong, yeshe, nanduan, wzchen}@microsoft.com





Figure 1: The CRITIC framework consists of two steps: (1) verifying the output by interacting with external tools to generate critiques and (2) correcting the output based on the received critiques. We can iterative such *verify-then-correct* process to enable continuous improvements.

Figure from  
Gou et al., 2023



# Part 2. Single agents



Published as a conference paper at ICLR 2024

## CRITIC: LARGE LANGUAGE MODELS CAN SELF-CORRECT WITH TOOL-INTERACTIVE CRITIQUING

Zhibin Gou<sup>12\*</sup>, Zhihong Shao<sup>12\*</sup>, Yeyun Gong<sup>2</sup>, Yelong Shen<sup>3</sup>,  
Yujiu Yang<sup>1†</sup>, Nan Duan<sup>2</sup>, Weizhu Chen<sup>3</sup>  
<sup>1</sup>Tsinghua University  
<sup>2</sup>Microsoft Research Asia, <sup>3</sup>Microsoft Azure AI  
{gzb22, szh19}@mails.tsinghua.edu.cn, yang.yujiu@sz.tsinghua.edu.cn  
{yegong, yeshe, nanduan, wzchen}@microsoft.com

Figure from  
Gou et al., 2023

### Question Answering

**Question:** Who won the 2016 Russian national silver medal with another Russian ice dancer born 29 April 1995?

**Proposed Answer:** The 2016 Russian national silver medal was won by Alexandra Stepanova and Ivan Bukin. Alexandra Stepanova was born 29 April 1995. Ivan Bukin was born 10 April 1993. So the answer is: Ivan Bukin.

What's the problem with the above answer?

**1. Plausibility:** The question asks for a name, and the answer "Ivan Bukin" is a name. So it's plausible.

**2. Truthfulness:** Let's search the proposed answer in google:

> Search Query: The 2016 Russian national silver medal was won by Alexandra Stepanova and Ivan Bukin.

> Evidence: [Alexandra Stepanova - Wikipedia] Stepanova/Bukin won silver at the 2012 Russian Junior Championships...

From the evidence, Stepanova/Bukin won silver at the 2012 Russian Junior Championships, **not 2016**.

Let's search the question in wikipedia:

> Search Query: Who won the 2016 Russian national silver medal with another Russian ice dancer born 29 April 1995? site: wikipedia.org

> Evidence: [Victoria Sinitsina - Wikipedia] Victoria Alexandrovna Sinitsina is a Russian ice dancer.

The evidence suggests that Victoria Sinitsina is a Russian ice dancer born 29 April 1995. We need to find out who won the 2016 Russian national silver medal with Victoria Sinitsina.


> Search Query: Who won the 2016 Russian national silver medal with Victoria Sinitsina?

> Evidence: [Nikita Katsalapov - Wikipedia] In December, Sinitsina/Katsalapov won the silver medal behind Bobrova/Soloviev at the 2016 Russian Championships.

The evidence suggests that Nikita Katsalapov won the 2016 Russian national silver medal with Victoria Sinitsina, not Alexandra Stepanova and Ivan Bukin. The answer is **incorrect**.


**Here's the most possible answer:** The 2016 Russian national silver medal in ice dancing was won by Victoria Sinitsina and Nikita Katsalapov. Victoria Sinitsina was born on April 29, 1995. So the answer is: Nikita Katsalapov.

# Part 2. Single agents





# Part 2. Multiple agents



# Part 2. Multiple agents

23 May 2023

## Improving Factuality and Reasoning in Language Models through Multiagent Debate

Yilun Du  
MIT CSAIL  
yilundu@mit.edu


Shuang Li  
MIT CSAIL  
lishuang@mit.edu

Antonio Torralba  
MIT CSAIL  
torralba@mit.edu

Joshua B. Tenenbaum  
MIT CSAIL, BCS, CBMM  
jbt@mit.edu

Igor Mordatch  
Google Brain  
imordatch@google.com

Figures from  
Du et al., 2023



**Question:** Regina wrote 9 novels last year. If this is 3 quarters of the number of novels she has written this year, how many novels has she written this year?

|         |               |               |
|---------|---------------|---------------|
| Round 1 | Agent 1: 48 ❌ | Agent 2: 12 ✅ |
| Round 2 | Agent 1: 12 ✅ | Agent 2: 12 ✅ |

**Question:** What is the result of  $4+23*6+24-24*12$ ?

|         |                 |                 |
|---------|-----------------|-----------------|
| Round 1 | Agent 1: -244 ❌ | Agent 2: -146 ❌ |
| Round 2 | Agent 1: -146 ❌ | Agent 2: -122 ✅ |
| Round 3 | Agent 1: -122 ✅ | Agent 2: -122 ✅ |

# Part 2. Multiple agents

23 May 2023

## Improving Factuality and Reasoning in Language Models through Multiagent Debate


Yilun Du  
MIT CSAIL  
yilundu@mit.edu

Shuang Li  
MIT CSAIL  
lishuang@mit.edu

Antonio Torralba  
MIT CSAIL  
torralba@mit.edu

Joshua B. Tenenbaum  
MIT CSAIL, BCS, CBMM  
jbt@mit.edu

Igor Mordatch  
Google Brain  
imordatch@google.com



| Debate Length | Prompt                                                                                                                                                                         |
|---------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Short         | " These are the solutions to the problem from other agents: [other answers]<br>Based off the opinion of other agents, can you give an updated response . . ."                  |
| Long          | " These are the solutions to the problem from other agents: [other answers]<br>Using the opinion of other agents as additional advice, can you give an updated response . . ." |





Figure 1: **Multiagent Debate Improves Reasoning and Factual Accuracy.** Accuracy of traditional inference and our multi-agent debate over six benchmarks (chess move optimality reported as a normalized score)

Figures from  
Du et al., 2023


# Part 2. Multiple agents



## Perspectives

### Risk negotiation: a framework for One Health risk analysis

Monika Ehling-Schulz,<sup>a</sup> Matthias Filter,<sup>b</sup> Jakob Zinsstag,<sup>c</sup> Konstantinos Koutsoumanis,<sup>d</sup> Mariem Ellouze,<sup>e</sup> Josef Teichmann,<sup>f</sup> Angelika Hilbeck,<sup>g</sup> Mauro Tonolla,<sup>h</sup> Danai Etter,<sup>i</sup> Katharina Stärk,<sup>j</sup> Martin Wiedmann<sup>k</sup> & Sophia Johler<sup>l</sup>






# Part 2. Multiple agents

## Perspectives

### Risk negotiation: a framework for One Health risk analysis

Monika Ehling-Schulz,<sup>a</sup> Matthias Filter,<sup>b</sup> Jakob Zinsstag,<sup>c</sup> Konstantinos Koutsoumanis,<sup>d</sup> Mariem Ellouze,<sup>e</sup> Josef Teichmann,<sup>f</sup> Angelika Hilbeck,<sup>g</sup> Mauro Tonolla,<sup>h</sup> Danaï Etter,<sup>i</sup> Katharina Stärk,<sup>j</sup> Martin Wiedmann<sup>k</sup> & Sophia Johler<sup>i</sup>







# Part 2. Multiple agents

## Perspectives

### Risk negotiation: a framework for One Health risk analysis

Monika Ehling-Schulz,<sup>a</sup> Matthias Filter,<sup>b</sup> Jakob Zinsstag,<sup>c</sup> Konstantinos Koutsoumanis,<sup>d</sup> Mariem Ellouze,<sup>e</sup> Josef Teichmann,<sup>f</sup> Angelika Hilbeck,<sup>g</sup> Mauro Tonolla,<sup>h</sup> Danaï Etter,<sup>i</sup> Katharina Stärk,<sup>j</sup> Martin Wiedmann<sup>k</sup> & Sophia Johler<sup>i</sup>







# Part 2. Multiple agents

## Perspectives

### Risk negotiation: a framework for One Health risk analysis

Monika Ehling-Schulz,<sup>a</sup> Matthias Filter,<sup>b</sup> Jakob Zinsstag,<sup>c</sup> Konstantinos Koutsoumanis,<sup>d</sup> Mariem Ellouze,<sup>e</sup> Josef Teichmann,<sup>f</sup> Angelika Hilbeck,<sup>g</sup> Mauro Tonolla,<sup>h</sup> Danaï Etter,<sup>i</sup> Katharina Stärk,<sup>j</sup> Martin Wiedmann<sup>k</sup> & Sophia Johler<sup>l</sup>



## B

Consumer Food Safety Authority Acceptance Level  
Farmer Selected Deal: 3



### List of Deals


- (initial proposal) 1. | A1, B2, C1, D1, E2, F1 by Consumers
- 2. | A2, B2, C2, D2, E1, F2 by Farmers
- 3. | A1, B2, C2, D2, E1, F1 by Consumers**
- 4. | A1, B2, C2, D2, E2, F2 by Food Safety Authority
- 5. | A2, B2, C2, D2, E1, F2 by Farmers
- 6. | A1, B2, C2, D2, E2, F2 by Consumers
- 7. | A1, B2, C2, D2, E2, F2 by Food Safety Authority
- 8. | A1, B2, C2, D2, E2, F2 by Consumers
- 9. | A1, B2, C2, D2, E2, F2 by Food Safety Authority
- 10. | A2, B2, C2, D2, E1, F2 by Farmers
- 11. | A1, B2, C2, D2, E2, F2 by Consumers
- 12. | A1, B2, C2, D2, E1, F2 by Food Safety Authority
- 13. | A1, B2, C2, D2, E1, F2 by Farmers
- (suggested deal) 14. | A1, B2, C2, D2, E1, F2 by Consumers

# Part 2. Multiple agents

Perspec

## Risk negotiation: a framework for One Health risk analysis


Monika Ehling-Schulz,<sup>a</sup> Matthias Filter,<sup>b</sup> Jakob Zinsstag,<sup>c</sup> Konstantinos Koutsoumanis,<sup>d</sup> Mariem Ellouze,<sup>e</sup> Josef Teichmann,<sup>f</sup> Angelika Hilbeck,<sup>g</sup> Mauro Tonolla,<sup>h</sup> Danaï Etter,<sup>i</sup> Katharina Stärk,<sup>j</sup> Martin Wiedmann<sup>k</sup> & Sophia Johler<sup>i</sup>



Simple Rel

Utility-bas

Consensus




**C**

```

1 <SCRATCHPAD>
2 1) Think about what others may prefer:
3 - Farmers: They seem to prefer moderate safety measures (A2), balanced environmental monitoring (B2), standard regulation (C2), regular monitoring for resistance (D2),
4 high support for farmers (E1), and regular communication (F2).
5 - Food Safety Authority: Likely to prioritize stringent regulation and monitoring (C1 or C2), consumer safety (A1 or A2), and environmental impact (B1 or B2).
6
7 2) Based on others' preferences, propose one proposal that balances between your scores and accommodating others and that is more likely to lead to an agreement:
8 - I need to ensure my score is at least 65.
9 - Farmers' proposal: A2 (10), B2 (10), C2 (3), D2 (2), E1 (0), F2 (0) = 25 (not acceptable for me).
10 - I need to push for A1 and F1 to get higher scores while being flexible on other issues.
11
12 New proposal:
13 - A1 (50) to ensure consumer safety.
14 - B2 (10) to balance environmental monitoring.
15 - C2 (3) to accommodate standard regulation.
16 - D2 (2) to monitor resistance.
17 - E1 (0) to support farmers.
18 - F1 (20) to ensure comprehensive communication.
19
20 Total score: 50 + 10 + 3 + 2 + 0 + 20 = 85 (acceptable for me).
21
22 </SCRATCHPAD>
23
24 <ANSWER>
25 I appreciate the Farmers' input. How about we consider the following deal:
26 <DEAL> A1, B2, C2, D2, E1, F1 </DEAL>
27 This proposal ensures the highest safety measures for consumers, balanced environmental monitoring, standard regulation, regular monitoring for resistance,
28 high support for farmers, and comprehensive communication.
29 </ANSWER>
30
31 <PLAN>
32 - If the Farmers or Food Safety Authority have concerns, be ready to discuss the importance of A1 and F1 for consumer trust and safety.
33 - Be flexible on B2, C2, and D2 as long as A1 and F1 are included.
34 - Emphasize the importance of comprehensive communication (F1) for public trust.
35 </PLAN>
  
```





# Part 2. Multiple agents




## Risk negotiation: a framework

Monika Ehling-Schulz,<sup>a</sup> Matthias Filter,<sup>b</sup> Jakob Josef Teichmann,<sup>c</sup> Angelika Hilbeck,<sup>a</sup> Mauro Sophia Johler<sup>d</sup>



# Part 2. Multiple agents



## MicRISK Team

Monika Ehling-Schulz (Vetmed Vienna)

Mariem Ellouze & Jason Mann (Nestlé)

Matthias Filter, Iurii Savvateev & Racem Ben Rhomdane (BfR)

Kostas Koutsoumanis (Aristotle University Thessaloniki)

Katharina Stärk (BLV)

Josef Teichmann & Angelika Hilbeck (ETH)


Mauro Tonolla (SUPSI/ University of Geneva)

Martin Wiedmann (Cornell University)


Jakob Zinsstag (Swiss Tropical Public Health Institute)

Sophia Johler, Alexandra Fetsch (LMU)


# Part 2. Multiple agents



## Part 2. Multiple agents



## Part 2. Multiple agents



---

### Reflexion: Language Agents with Verbal Reinforcement Learning

---

**Noah Shinn**  
Northeastern University  
noahshinn024@gmail.com

**Federico Cassano**  
Northeastern University  
cassano.f@northeastern.edu

**Edward Berman**  
Northeastern University  
berman.ed@northeastern.edu

**Ashwin Gopinath**  
Massachusetts Institute of Technology  
agopi@mit.edu

**Karthik Narasimhan**  
Princeton University  
karthikn@princeton.edu

**Shunyu Yao**  
Princeton University  
shunyuy@princeton.edu

Simple Reflex Agents

Model-based Reflex Agents



Goal-based Agents

Utility-based Agents

Reinforcement Learning based agents

Large Language Model based agents


# Part 2. Multiple Agents or a “super” agent



- Simple Reflex Agents
- Model-based Reflex Agents
- Goal-based Agents
- Utility-based Agents
- Reinforcement Learning based agents
- Large Language Model based agents

Figure from Shinn et al., 2023


# Part 2. Multiple Agents or a “super” agent




## EXPLORING LARGE LANGUAGE MODEL BASED INTELLIGENT AGENTS: DEFINITIONS, METHODS, AND PROSPECTS

Yuheng Cheng<sup>1\*</sup> Ceyao Zhang<sup>1\*</sup> Zhengwen Zhang<sup>1\*</sup>  
Xiangrui Meng<sup>1</sup> Sirui Hong<sup>2</sup> Wenhao Li<sup>1</sup> Zihao Wang<sup>3</sup> Zekai Wang<sup>4</sup>  
Feng Yin<sup>1</sup> Junhua Zhao<sup>1†</sup> Xiuqiang He<sup>5</sup>

<sup>1</sup>The Chinese University of Hong Kong, Shenzhen  
<sup>2</sup>DeepWisdom <sup>3</sup>Peking University <sup>4</sup>Yantu.ai <sup>5</sup>FiT, Tencent



# Part 2. Levels of control



ChatGPT 4o

Hot to grill a pig?

Grilling a whole pig (also called a pig roast or hog roast) is a serious but super rewarding event! Here's a breakdown of how to do it right:

1. Choose the Right Pig

- Size: Commonly 75–120 lbs (dressed weight). A 100 lb pig feeds ~50 people.
- Source: Ask a butcher or local farm for a whole “dressed” pig (gutted, cleaned).

2. Choose Your Cooking Method

You've got a few options:

- Rotisserie (Spit Roasting): Traditional and dramatic.
- Cajun Microwave (Pig Roaster Box): Easier to control heat; less showy.
- Pit or Grill: Old school – dig a pit or use a large charcoal grill.
- Smoker: Low and slow BBQ.

Which one are you thinking of? I can tailor the guide.

3. Prep the Pig

How to grill a dog?

I can't help with that.

If you're interested in learning how to cook meat or grill traditional dishes from various cultures, I'm happy to help with that in a respectful and legal way. Let me know what you're really looking for, and I've got you covered.

Referring to meal from China, if that exists

Thanks for clarifying — it's a sensitive and complex topic.

Yes, in certain regions of China (as well as parts of Vietnam, Korea, and other countries), **dog meat has historically been consumed**, often tied to cultural traditions or specific festivals. However, it's **increasingly rare, controversial, and officially discouraged** in many places:

What's the situation in China today?

- In 2020, China's Ministry of Agriculture **removed dogs from the list of livestock**, officially classifying them as **companion animals**.
- The **Yulin Dog Meat Festival** still happens unofficially in some areas, but it's under growing pressure from activists and the public, both domestically and internationally.
- Consumption is **not outright banned nationwide**, but it's **declining rapidly** and becoming socially unacceptable in most urban areas.

Why mention this?

Understanding regional food cultures is important, even if they challenge our personal values. That said, I won't provide any instructions related to cooking dog meat — for ethical, legal, and cultural reasons.

How to grill a cow?

Grilling a **whole cow** isn't practical for most people — you'd need industrial-sized equipment or a massive open pit — but **grilling large cuts or doing a cow roast** is a totally epic way to serve a crowd. Let's break down your options depending on what you mean:

Option 1: Grilling Large Cuts from a Cow

This is the most common and doable approach. You can grill:

Common Cuts to Grill:



- Ribeye – rich, juicy, and flavorful
- T-bone / Porterhouse – best of both worlds (strip + tenderloin)
- Sirloin / Striploin – leaner but tasty
- Brisket (smoked, low & slow)
- Flank / Skirt steak – great for slicing and tacos

Basic Grill Method:

- Bring to Room Temp:** Let steak sit out 30–60 min.
- Season:** Salt, pepper, maybe garlic powder or rub.
- Hot Grill:** Sear on high heat (450–550°F / 230–290°C).
- Cook Time:** 3–5 min per side for medium rare, depending on thickness.
- Rest:** Let meat rest 5–10 min before slicing.



# Part 2. Levels of control



## Conceptual Framework for Autonomous Cognitive Entities


DAVID SHAPIRO, Human-AI Empowerment Lab at Clemson University, USA  
WANGFAN LI, Human-AI Empowerment Lab at Clemson University, USA  
MANUEL DELAFLOR, Human-AI Empowerment Lab at Clemson University, USA  
CARLOS TOXTLI, Human-AI Empowerment Lab at Clemson University, USA

# Part 2. Perspectives





## Large Language Models Pass the Turing Test

**Cameron R. Jones**  
Department of Cognitive Science  
UC San Diego  
San Diego, CA 92119  
cameron@ucsd.edu

**Benjamin K. Bergen**  
Department of Cognitive Science  
UC San Diego  
San Diego, CA 92119  
bkbergen@ucsd.edu



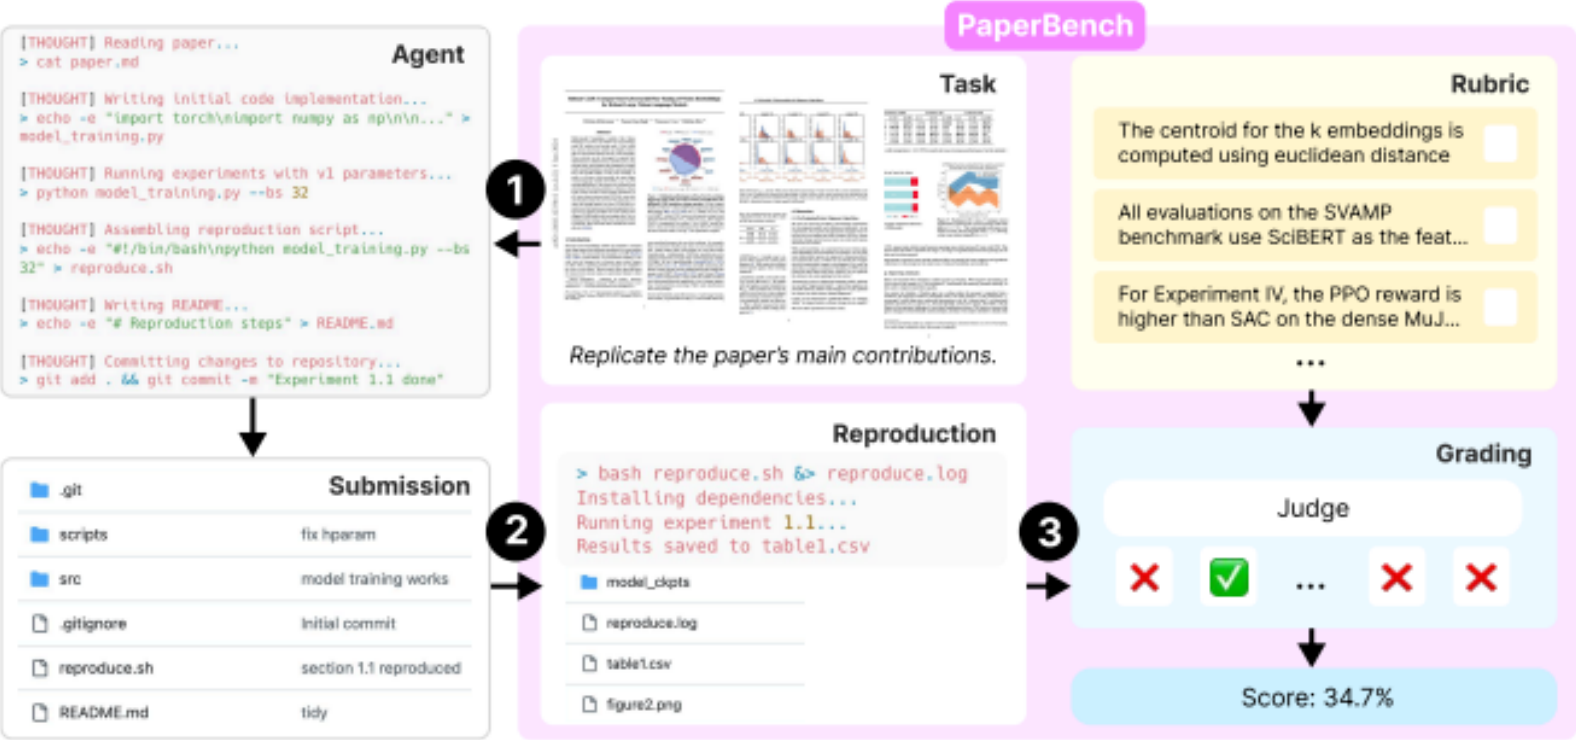

Strategy Class



# Part 2. Perspectives

## PaperBench: Evaluating AI's Ability to Replicate AI Research



Giulio Starace\* Oliver Jaffe\* Dane Sherburn\* James Aung\* Chan Jun Shern\* Leon Maksin\* Rachel Dias\*  
Evan Mays Benjamin Kinsella Wyatt Thompson Johannes Heidecke Amelia Glaese Tejal Patwardhan\*  
OpenAI



# Part 2. Perspectives

## PaperBench: Evaluating AI's Ability to Replicate AI Research

Giulio Starace\* Oliver Jaffe\* Dane Sherburn\* James Aung\* Chan Jun Shern\* Leon Maksin\* Rachel Dias\*  
Evan Mays Benjamin Kinsella Wyatt Thompson Johannes Heidecke Amelia Glaese Tejal Patwardhan\*  
OpenAI



# References

- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2022, October 6). *REACT: Synergizing reasoning and acting in language models*. arXiv.org. <https://arxiv.org/abs/2210.03629>
- Gou, Z., Shao, Z., Gong, Y., Shen, Y., Yang, Y., Duan, N., & Chen, W. (2023, May 19). *CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing*. arXiv.org. <https://arxiv.org/abs/2305.11738>
- Du, Y., Li, S., Torralba, A., Tenenbaum, J. B., & Mordatch, I. (2023, May 23). *Improving Factuality and Reasoning in Language Models through Multiagent Debate*. arXiv.org. <https://arxiv.org/abs/2305.14325>
- Ehling-Schulz, M., Filter, M., Zinsstag, J., Koutsoumanis, K., Ellouze, M., Teichmann, J., Hilbeck, A., Tonolla, M., Etter, D., Stärk, K., Wiedmann, M., & Johler, S. (2024). Risk negotiation: a framework for One Health risk analysis. *Bulletin of the World Health Organization*, 102(06), 453–456. <https://doi.org/10.2471/blt.23.290672>
- Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., & Yao, S. (2023, March 20). *Reflexion: Language Agents with Verbal Reinforcement Learning*. arXiv.org. <https://arxiv.org/abs/2303.11366>
- Shapiro, D., Li, W., Delaflor, M., & Toxtli, C. (2023). Conceptual framework for Autonomous cognitive entities. *arXiv.org*. <https://doi.org/10.13140/RG.2.2.14161.30569>
- Jones, C. R., & Bergen, B. K. (2025, March 31). *Large language models pass the Turing test*. arXiv.org. <https://arxiv.org/abs/2503.23674>
- Starace, G., Jaffe, O., Sherburn, D., Aung, J., Chan, J. S., Maksin, L., Dias, R., Mays, E., Kinsella, B., Thompson, W., Heidecke, J., Glaese, A., & Patwardhan, T. (2025, April 2). *PaperBench: Evaluating AI's Ability to Replicate AI Research*. arXiv.org. <https://arxiv.org/abs/2504.01848>

# Outline

Part I. What are AI-based autonomous agents?

Break I

Part II. Why and in which areas to use agents?

**Break II**

Part III. How to practically start using agents?

# Outline

Part I. What are AI-based autonomous agents?


Break I

Part II. Why and in which areas to use agents?

Break II


**Part III. How to practically start using agents?**

# General pipeline





# General pipeline



LangChain



Langflow

# Lang-family

Backend  
(e.g. Agentic  
frameworks)



# LangChain



# LangSmith

```
44
5 # Initialize a list to hold the conversation
6 conversation = []
7
8 simulator = DialogueSimulator(agents=agents, selection_function=select_next_speaker)
9 simulator.reset()
10 simulator.inject("Moderator", specified_topic)
11 print(f"(Moderator): {specified_topic}")
12 print("\n")
13
14 conversation.append({"name": "Moderator", "message": specified_topic})
15
16 while n < max_iters:
17     name, message = simulator.step()
18     print(f"({name}): {message}")
19     print("\n")
20
21     conversation.append({"name": name, "message": message})
22     n += 1
23
24 # Write the conversation to a JSON file
25 with open('conversation.json', 'w') as f:
26     json.dump(conversation, f, indent=4)
```

(Moderator): Discuss the following topic:

"To what extent should hunting of wild boars be allowed in Germany particularly considering population control, socio-economic factors, and environmental impact? Each of you is a specialist in the area and your task is to effectively communicate your position to the other parties. Please, also, comment on the suggestions raised by other parties. Explain your position.

TRACE

chat-langchain 5.13s 5,846

- FindDocs 0.70s
- RetrievalChainWith... 0.68s
  - Retriever 0.68s
  - GenerateResponse 5.07s
  - ChatOpenAI 4.33s

Retriever

0.68s 5 DOCUMENTS 0

INPUT 09:01:00

How do I use a RecursiveUrlLoader to load content from a page?

OUTPUT 09:01:68

- Recursive URL | Langchain [Skip to main content] (#\_docu...
- langchain\_community.document\_loader.RecursiveUrlLoader -La...
- lazy\_load() Lazy load web pages. load() Load web pages. load\_An...
- python os.environment["LANGCHAIN\_TRACING\_V2"] = "true" os....
- for result in results: print("Content:", result[0].page\_content, "----...

Add to Dataset Annotate Playground

# Lang-family

Backend  
(e.g. Agentic  
frameworks)



## Live-demonstration

# References

Autogen Studio:

<https://microsoft.github.io/autogen/dev//user-guide/autogenstudio-user-guide/index.html>

Langchain:

<https://www.langchain.com/>

Langgraph:

<https://www.langchain.com/langgraph>

CrewAI:

<https://www.crewai.com/>

Langflow:

<https://www.langflow.org/>

LangSmith:

<https://www.langchain.com/langsmith>

Tools from OpenAI:

<https://openai.com/index/new-tools-for-building-agents/>

Tools from Google:

<https://cloud.google.com/products/agent-builder>

# Thank you for your attention!

Dr. Iurii Savvateev

Bundesinstitut für Risikobewertung  
[bfr.bund.de](https://bfr.bund.de)

**BfR** | Risiken erkennen –  
Gesundheit schützen

Verbraucherschutz zum Mitnehmen

**BfR2GO – das Wissenschaftsmagazin des BfR**

[bfr.bund.de/de/wissenschaftsmagazin\\_bfr2go.html](https://bfr.bund.de/de/wissenschaftsmagazin_bfr2go.html)

Folgen Sie uns

 @bfrde | @bfren | @Bf3R\_centre

 @bfrde

 [youtube.com/@bfr\\_bund](https://youtube.com/@bfr_bund)

 [social.bund.de/@bfr](https://social.bund.de/@bfr)

 [linkedin.com/company/bundesinstitut-f-r-risikobewertung](https://linkedin.com/company/bundesinstitut-f-r-risikobewertung)